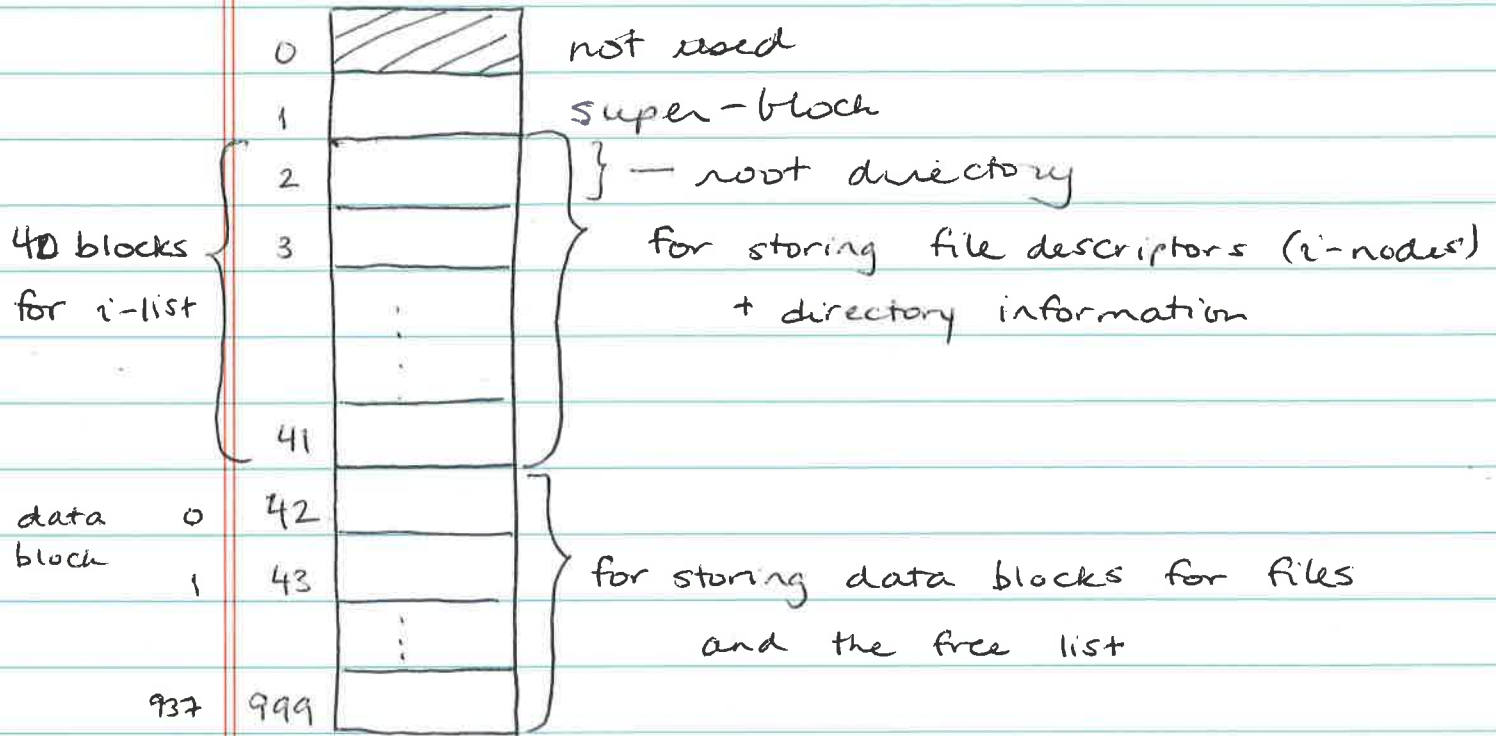


Unix V6 file system

1000 blocks of 512 bytes




relevant fields of an super block

i-size : # blocks in the i-list, here 40

n-free : index into free

free[100] :



if = 0, no more free blocks

Suppose there were only 20 data blocks
and that the free list was of length 5:

0	1, 4, 16, 17, 18, 19
1	2, 4, 12, 13, 14, 15
2	3, 4, 8, 9, 10, 11
3	4, 4, 4, 5, 6, 7
4	0, 4, 0, 1, 2, 3
5	
6	
7	
⋮	
⋮	
19	

} free list

1st entry = block contain next set
of 4 free blocks
= 0, then no more of them

2nd entry
= # free blocks in the list, n_{free}


list of n_{free} free blocks

so free[5]

	0	1	2	3	4
free	1	16	17	18	19

n_{free} = 4

how to allocate a file f



suppose 720 bytes in size
since each block is 512
bytes need 2 blocks

~~free~~ free[4] = 19, nfree = 4 - 1 = 3

copy b0 of f into block 19, free[3] = 18

copy b1 of f into block 18, nfree = 2

now allocate file g of size 1224 bytes
⇒ need 3 blocks

copy b0 of g into free[2] = block 17, nfree = 1

copy b1 of g into free[1] = block 16, nfree = 0

since nfree = 0 ⇒ block free[1] = 1

contains next free list

update

free

2	12	13	14	15
---	----	----	----	----

nfree = 4

now allocate ^{copy} b2 of g into free[4] = block 15
nfree = 3

i-nodes (file descriptors) + directories

i-numbers begin at 1; storage starts in block 2

i-nodes are 32 bytes long so. ($32 \times 16 = 512$)

16 of them fit in one block

root directory in block 2

each i-node represents one file

flags		flag on file	000700					
n links		# links to file	0					
uid		user id of owner	say user 1					
gid		group id of owner						
size 0		} 24 bits for file size	= 1224 bytes					
size 1								
addr	0	1	2	3	4	5	6	7
atime	0	1	time of last access					
mtime	0	1	time of last modification					

will only deal with small files ($8 \times 512 \text{b} = 4\text{K}$)

\Rightarrow so $\text{addr}[i]$, $0 \leq i \leq 7$ are direct address blocks, i.e., give block # of file contents


for previous example of file g

	0	1	2	3	4	5	6	7
addr	17	16	15					

directory entries

indicated by flag bit of i-node entry

directory entries are 16 bytes long \Rightarrow 32 in one block



if = 0 then entry is empty
o/w it is the file number

by convention 1st 2 entries in each directory are . and ..

current parent

in block 2 is the root directory

i.e., can contain 16 i-nodes

flag has 140777 inode r/w/x

addr array for a directory is

contents of the directory, i.e., if

addr[0] = 57 then block 57

contains directory entries (up to 32 of them)

	size	
0		.
1		..
2		/
3	!	!

- main thread

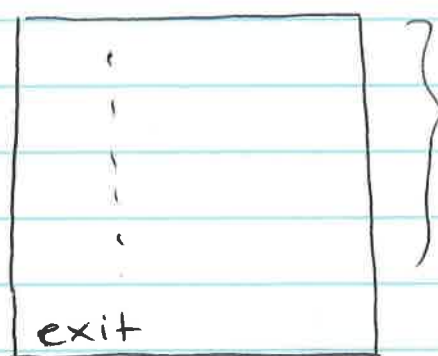
ubfs disk k

initializes the file system

spawns k threads

each thread i , $1 \leq i \leq k$

file user i .txt

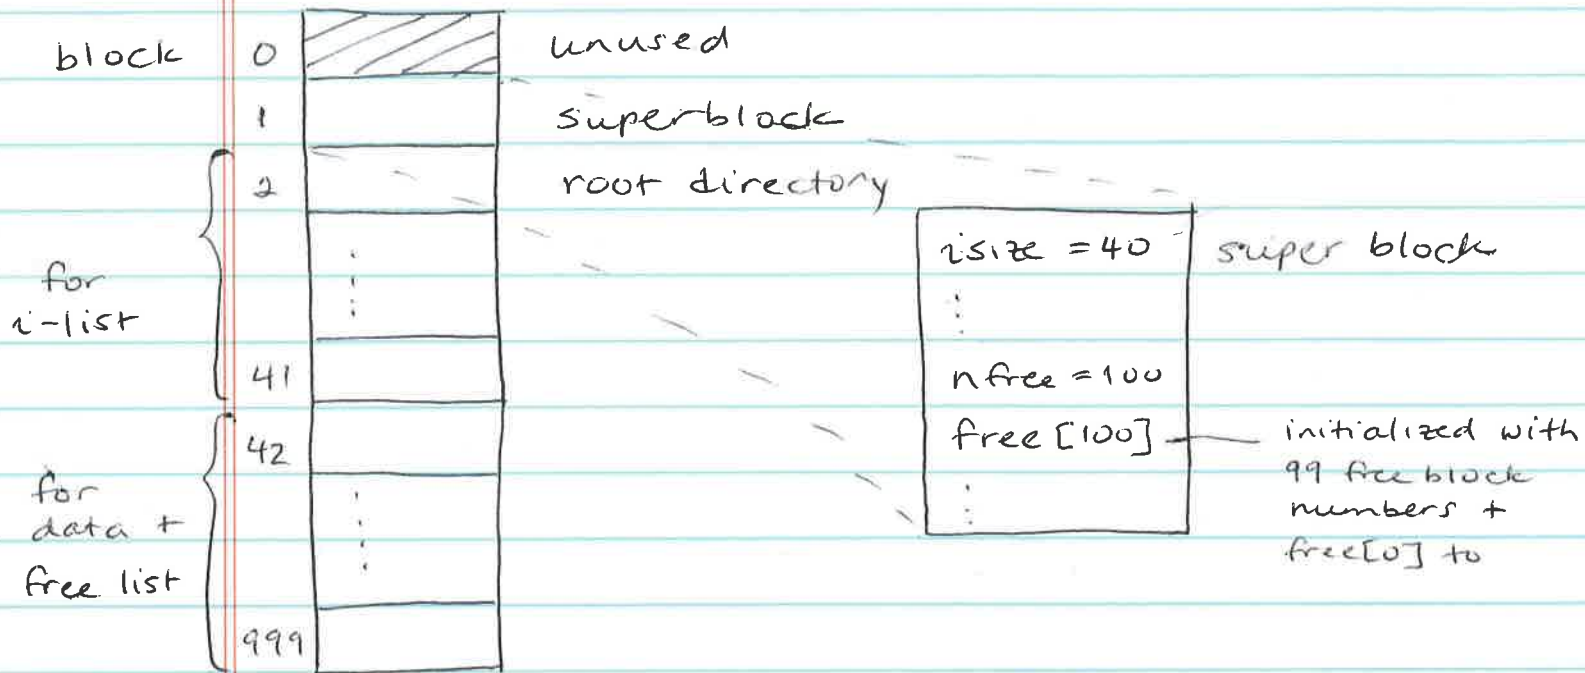


sequence of commands

- file system is allocated in memory

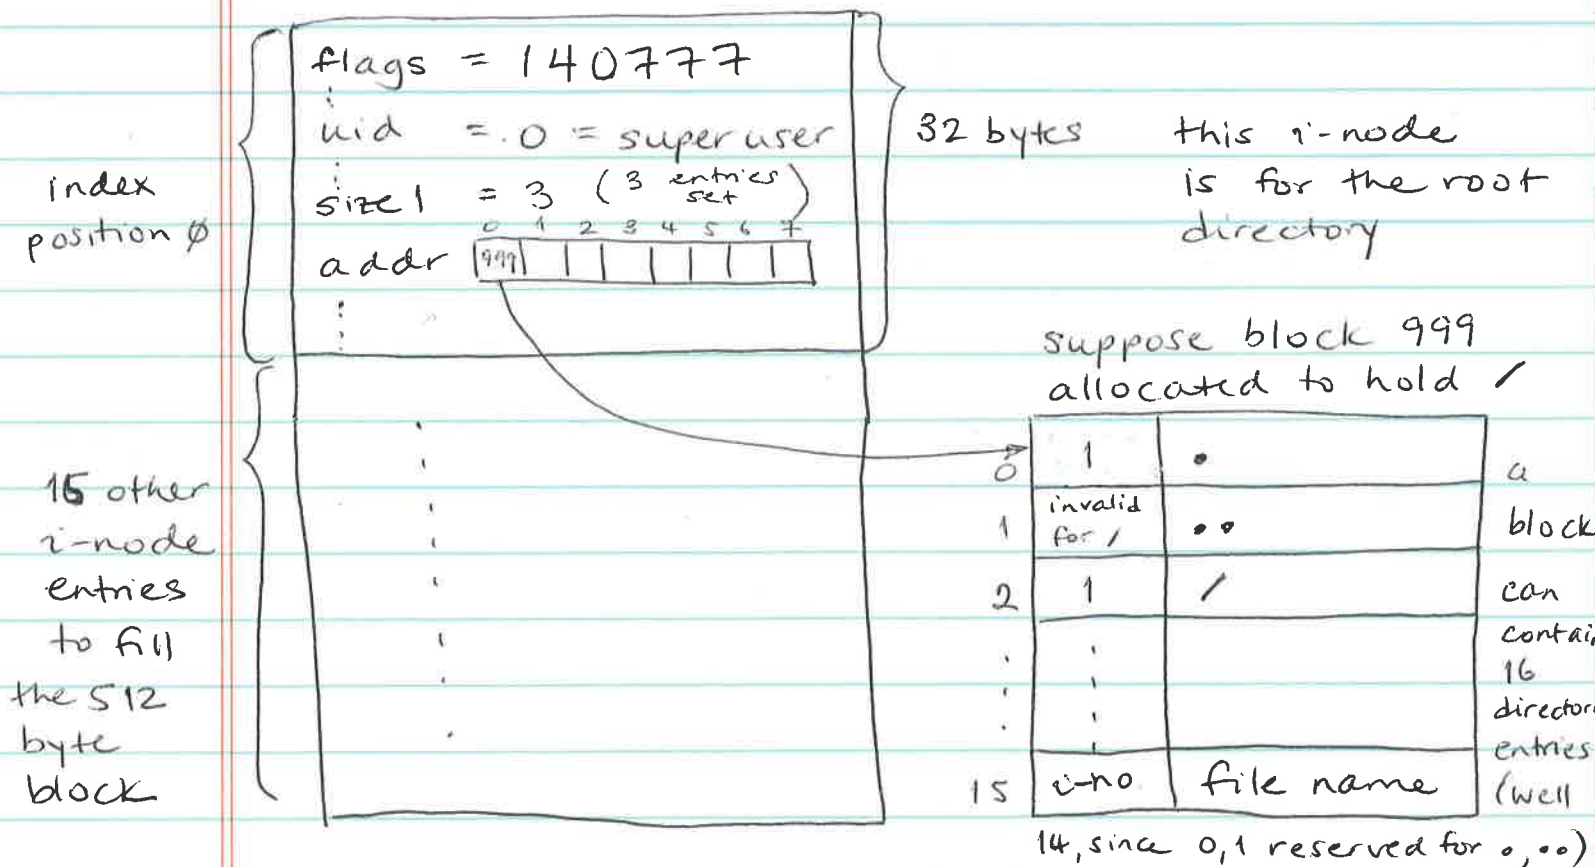
- work on it in memory

After file system is initialized it will contain



root directory (block 2)

overlay an array of 16 i-nodes onto block 2



First, suppose a user 1 creates a directory d1; since the current directory is initialized as /, d1 gets created in /

To create a directory:

- 1) Create an entry in the current directory for it

⇒ into block 999 (holding /)

block 999


also update size of / i-node to be 4 (for 4 entries are set)

0	1	.
1	-1	..
2	1	/
3	2	d1
	:	

assume i-number = -1 means invalid (can't go up from /)

- 2) Create an i-node for the directory d1
⇒ into block 2, add an i-node for d1 at position 1

block 2



Note that creating the directory allocated a block to hold entries for the directory (initialized with ., and ..)

Suppose user 1 now did `cd to d1`, and then executed `"cpin f1.txt f1v6.txt"`

The `cd` now has block 998 as the current directory. Hence the `cpin` will create the directory entry for the file here

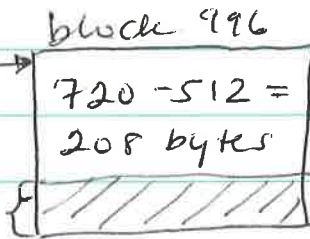
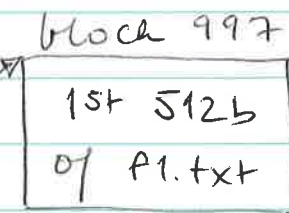
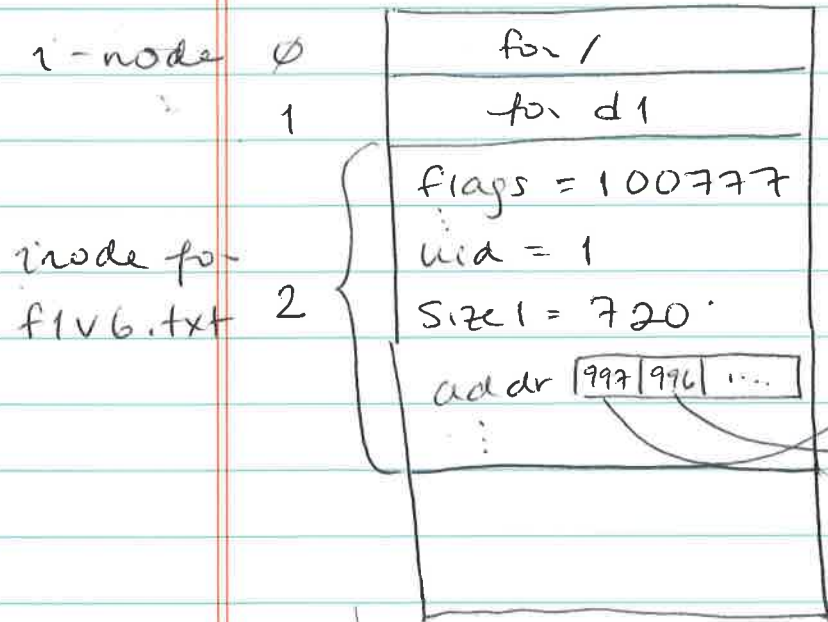
block 998 (directory d1)

also update size of d1 to 3

2	.
1	..
3	f1v6.txt
⋮	⋮

- 1) Just create a directory entry for the file
- 2) Create an i-node for the file (suppose f1.txt has 720 bytes)

block 2



internal fragmentation

If move back up to / and create 1 entries in block 999, would need to allocate another block to directory entries and update `addr[1]` in block 2 (1st i-node in block 2)

For directories `size1` = number of directory entries in the block

For files, `size1` = the file size in bytes; since our max. file size is 4K, the `size0` field is unused.